


## 9.1 Build a Docker Application with Kubernetes Service (Example 1)



### Objectives:

- Create a Kubernetes cluster.
- Build a Docker application with azure Kubernetes service.

### Part 1. Create a Kubernetes Cluster

1. From Aure Portal, create a new resource. Select Containers➔ AKS.

## Create a resource ...

Get Started

Recently created

### Categories

AI + Machine Learning

Analytics

Blockchain

Compute

Containers

Popular Azure services [See more in All services](#)



Azure Kubernetes Service (AKS)

[Create](#) | [Docs](#) | [MS Learn](#)



Web App for Containers

[Create](#) | [Docs](#) | [MS Learn](#)

















Batch Service

[Create](#) | [Docs](#) | [MS Learn](#)







2. Assign cluster a name and use default settings, including auto scale. Click “**Review + create**”.

**Cluster details**

Cluster preset configuration	Standard (\$\$)  <small>To quickly customize your Kubernetes cluster, choose one of the preset configurations above. You can modify these configurations at any time. <a href="#">Learn more and compare presets</a></small>
Kubernetes cluster name * 	azurebook 
Region * 	(US) East US 
Availability zones 	Zones 1,2,3   High availability is recommended for standard configuration.
AKS pricing tier 	Standard 
Kubernetes version * 	1.26.6 (default) 
Automatic upgrade 	Enabled with patch (recommended) 

**Primary node pool**

The number and size of nodes in the primary node pool in your cluster. For production workloads, at least 3 nodes are recommended for resiliency. For development or test workloads, only one node is required. If you would like to add additional node pools or to see additional configuration options for this node pool, go to the 'Node pools' tab above. You will be able to add additional node pools after creating your cluster.  
[Learn more about node pools in Azure Kubernetes Service](#)


Node size * 	<b>Standard DS2 v2</b>  Standard DS2_v2 is recommended for standard configuration. <a href="#">Change size</a>
Scale method * 	<input type="radio"/> Manual <input checked="" type="radio"/> Autoscale  Autoscaling is recommended for standard configuration.
Node count range * 	<input type="text" value="1"/>  <input type="text" value="5"/>

< Previous Next : Node pools > Review + create

3. Click “Create”.
4. Go to resource when deployment is complete.

## Part 2. Connect to the Cluster

5. Open Cloud Shell using the button on the top of the Azure portal.



6. Connect to your Kubernetes cluster using the `az aks get-credentials` command. Replace `myResourceGroup` and `myAKSCluster` with your own values.

```
$ az aks get-credentials --resource-group myResourceGroup --name myAKSCluster
```

7. In my case the command is the following.

```
$ az aks get-credentials --resource-group azurebook --name azurebook
```

8. Verify the connection to your cluster using the following command.

```
$ kubectl get nodes
```

9. This is a sample output. Make sure the node status is Ready.

NAME	STATUS	ROLES	AGE	VERSION
aks-agentpool-42033415-vmss000000	Ready	agent	22m	v1.26.6

### Part 3. Deploy the Application

10. In the Cloud Shell, upload file `azure-vote.yaml` given online or use the text editor to create this file.



11. Use command `ls` to show this file.

```
$ ls
```

12. This is the output.

```
azure-vote.yaml  clouddrive
```

13. Note, in file `azure-vote.yaml`, the front end image was created in previous chapters and is stored in Docker hub.

```
containers:  
- name: azure-vote-front  
  image: azurebook/voting_redis-vote_front
```

14. Deploy the application.

**\$ kubectl apply -f azure-vote.yaml**

15. Output is like the following.

```
deployment.apps/azure-vote-back created  
service/azure-vote-back created  
deployment.apps/azure-vote-front created  
service/azure-vote-front created
```

16. Use the following command to monitor the service.

**\$ kubectl get service azure-vote-front**

17. This is the output.

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
azure-vote-front	LoadBalancer	10.0.116.58	20.75.181.185	80:31147/TCP	2m42s

18. To test your application, open a web browser to the external IP address of your service shown at Step 17.

### Vote the Successor of Humankind

Deers

Dogs

Cats

Crows

Robots

Deers - 0 | Dogs - 1 | Cats - 0 | Crows - 0 | Robots - 0

## Part 4. Manually Scale Pods

19. When you deployed the application, a single replica was created for both the frontend and the backend services. To see the number and state of pods in your cluster, use the following command.

```
$ kubectl get pods
```

20. The output looks like the following.

NAME	READY	STATUS	RESTARTS	AGE
azure-vote-back-65c595548d-sxf22	1/1	Running	0	8m45s
azure-vote-front-65f64d9d4d-lrn46	1/1	Running	1 (8m12s ago)	8m44s

21. Manually change the number of pods in the azure-vote-front deployment using the **kubectl scale** command. The following example command increases the number of front-end pods to two.

```
$ kubectl scale --replicas=2 deployment/azure-vote-front
```

22. Verify the additional pods were created.

```
$ kubectl get pods
```

23. The output looks like the following.

azure-vote-back-65c595548d-sxf22	1/1	Running	0	34m
azure-vote-front-65f64d9d4d-5m6s4	1/1	Running	0	39s
azure-vote-front-65f64d9d4d-lrn46	1/1	Running	1 (34m ago)	34m

## Part 5. Auto Scale Pods

24. Kubernetes supports autoscaling to adjust the number of pods in a deployment depending on CPU utilization or other metrics.
25. To use the autoscaler, all containers and pods must have defined CPU requests and limits. In the azure-vote-front deployment, the front-end container requests 0.1 CPU with a limit of 0.25 CPU as shown below.

```
containers:
- name: azure-vote-front
  image: azurebook/voting_redis-vote_front
  resources:
    requests:
      cpu: 100m
      memory: 128Mi
    limits:
      cpu: 250m
      memory: 256Mi
```

26. The following command autoscales the number of pods in the azure-vote-front deployment with the following conditions: if average CPU utilization across all pods exceeds 50% of the requested usage, the autoscaler increases the pods up to a maximum of 5 instances and a minimum of 3 instances for the deployment.

```
$ kubectl autoscale deployment azure-vote-front --cpu-percent=50 --min=3 --max=5
```

## **Part 6. Scale the Cluster Nodes**

27. Get the number of nodes using the **az aks show** command.

```
$ az aks show --resource-group azurebook --name azurebook --query agentPoolProfiles
```

28. It shows the following info: one node is deployed. It can be scaled between 1 and 5.

```
"count": 1,
"creationData": null,
"currentOrchestratorVersion": "1.26.6",
"enableAutoScaling": true,
"enableEncryptionAtHost": null,
"enableFips": false,
"enableNodePublicIp": false,
"enableUltraSsd": null,
"gpuInstanceProfile": null,
"hostGroupId": null,
"kubeletConfig": null,
"kubeletDiskType": "OS",
"linuxOsConfig": null,
"maxCount": 5,
"maxPods": 110,
"minCount": 1,
"mode": "System",
"name": "agentpool",
```

29. Because auto scale is selected when we create azure Kubernetes service, we cannot manually scale the cluster.
30. If manual is selected when we create azure Kubernetes service, we can scale our cluster nodes using the **az aks scale** command.

```
$ az aks scale --resource-group azurebook --name azurebook --node-count 3
```

## **Part 8. Delete Service and Restart the Services**

31. Use the following command to list all services.

```
$ kubectl get service
```

32. The output is like following.

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
azure-vote-back	ClusterIP	10.0.130.237	<none>	6379/TCP	3m8s
azure-vote-front	LoadBalancer	10.0.138.59	52.150.42.5	80:30649/TCP	3m8s
kubernetes	ClusterIP	10.0.0.1	<none>	443/TCP	74m

33. Use the following command to delete the front service.

```
$ kubectl delete service azure-vote-front
```

34. Use the following command to delete the backend service.

```
$ kubectl delete service azure-vote-back
```

35. List all services.

```
$ kubectl get service
```

36. The output is like the following.

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.0.0.1	<none>	443/TCP	77m

37. Use the following command to restart the application.

```
$ kubectl apply -f azure-vote.yaml
```

38. Test the application in a web browser.
39. Delete all resources from Azure Portal.